

Designing a Versatile Multi-Agent AI Dialogue System

Overview: Multi-Agent Collaboration Benefits

Multi-agent AI systems involve multiple specialized **agents** (LLM instances with distinct roles or knowledge) working together, often outperforming a single monolithic AI on complex tasks ¹. By **leveraging multiple perspectives and feedback loops among agents**, these systems can tackle problems requiring reasoning, creativity, or multi-domain expertise more effectively than a single model alone ¹. Recent research shows that having agents critique and build on each other's outputs leads to more accurate and well-reasoned solutions, as each agent can catch others' mistakes and refine their answers ² ³. In fact, multi-agent approaches have been applied to *real-world simulations* in areas like psychology, policy making, and collaborative decision-making, illustrating their ability to model complex social or organizational environments ⁴. This makes them well-suited for use cases such as AI-driven debates or virtual advisory panels.

Key advantages of multi-agent systems include: breaking down a complex problem into specialized sub-tasks, cross-verification of facts, diverse brainstorming, and **consensus-building** across differing viewpoints ⁵. For example, an ensemble of agents can collectively draft and critique a plan – akin to a “team of experts” – yielding a more robust outcome than one AI working in isolation. This collaborative intelligence does come with trade-offs: coordination overhead and computational cost are higher, and new challenges (like managing agent interactions and ensuring consistency) must be addressed ⁶. The system design, therefore, needs to be **versatile** – supporting various interaction patterns – and robust in managing the complexity of multi-agent dialogues.

Multi-Agent Orchestration Patterns

To support a wide range of scenarios, the system should implement several **orchestration patterns** that define how agents coordinate and communicate ⁷ ⁸. Common patterns include **Sequential pipelines**, **Concurrent (parallel) agents**, and **Group chat (interactive debate)** – each suited to different problem structures:

- **Sequential Orchestration (Pipeline):** Agents operate in a fixed **linear sequence**, where each agent's output feeds into the next ⁹. This resembles a pipeline of specialists: e.g. one agent generates a draft, the next agent edits it, a third agent fact-checks, and so on. Such a chain **builds on dependencies step-by-step**, ideal for workflows with clear stages or progressive refinement (draft → review → polish) ¹⁰ ¹¹. The sequence is predetermined by the workflow; agents do not deviate or loop back upstream ¹⁰. For instance, a document generation pipeline might use an agent to choose a template, then another to fill in content, then a compliance agent to check regulations, each stage handing off to the next ¹² ¹³. This pattern ensures a structured handoff but is **less flexible** if iteration or backtracking is required ¹⁴.
- **Concurrent Orchestration (Parallel Agents):** Multiple agents run **simultaneously** on the same task, each providing independent analysis from a unique angle ¹⁵. The idea is to *fan-out* a problem to different experts in parallel and then aggregate their findings. This is useful when a

problem benefits from **diverse viewpoints or skill sets applied in parallel**, or when time is of the essence ¹⁶ ¹⁷. For example, given a question, one agent could reason logically, another could analyze quantitatively, and a third could bring creative insight – all at once. Their results might then be combined or voted on to produce a final answer ² ¹⁸. Parallel agents **do not directly talk to each other during processing**; instead, a coordinator collects their outputs and reconciles them (through majority vote, averaging, or a selection heuristic) ¹⁸ ¹⁹. This pattern accelerates processing and broadens coverage of the solution space, though it requires a clear strategy to handle possibly **conflicting results** from agents ²⁰.

- **Group Chat Orchestration (Interactive Debate/Collaboration):** Here multiple agents **actively converse in a shared thread**, interacting in turns to solve a problem or reach a decision ²¹. A dedicated *chat manager* (or moderator agent) oversees the dialogue, deciding which agent speaks when and enforcing the rules of engagement ²². All agent messages accumulate in a common conversation context that each participant can read and respond to ²³ ²⁴. This pattern shines when a task is best addressed via **brainstorming, debate, or iterative peer review**, much like a human committee meeting ²⁵. It supports free-flowing idea generation and **consensus-building through discussion** ²⁵. For example, agents with different expertise can propose solutions, critique each other, defend their points, and gradually converge on an agreed plan ²⁶. This is ideal for **creative ideation, complex decision-making** that benefits from arguing pros and cons, and **multidisciplinary problems requiring cross-functional dialogue** ²⁷. Unlike sequential pipelines, group-chat agents collaborate in real-time rather than simply handing off work ²⁸.

Each pattern can be instantiated depending on the use case. The **architecture should be flexible** enough to allow switching or blending patterns. For instance, a system might run several agents in parallel to generate ideas (concurrent brainstorming) and then have those agents enter a structured discussion to evaluate the ideas (group chat debate), followed by a sequential refinement pipeline to polish the chosen solution. The design must also handle **dynamic orchestration** – sometimes called the *handoff* or *delegate* pattern – where agents can **route tasks to whichever agent is best suited** on the fly ²⁹. In dynamic delegation, an agent assesses a query and, if it falls outside its expertise, passes it to a more appropriate agent (or even *spins up* a new specialist agent) ²⁹. This ensures **versatility**: the system can incorporate new experts as needed and scale to complex, heterogeneous tasks. A prominent example is MetaGPT's approach, which assigns distinct roles (e.g. product manager, architect, engineer) to different GPT agents and orchestrates their collaboration through a coordination mechanism ³⁰. MetaGPT's framework even allows recruiting additional specialized agents for specific needs, demonstrating how dynamic role allocation can tackle intricate tasks flexibly ²⁹.

Coordinating Agent Interactions

Regardless of pattern, a critical component is the **coordination logic** that manages agent interactions. In a sequential chain, this might be a simple controller that triggers each agent in order with the prior agent's output. For parallel agents, an aggregator component dispatches the task to all agents and later collects and integrates their responses (for example, via majority voting or by selecting the most confident answer) ² ¹⁸. In group chat mode, a more sophisticated **Chat Manager** is needed ²². The chat manager governs the conversation: it may enforce turn-taking (e.g. round-robin or a specific sequence like *maker* → *checker* → *maker* in a critique loop ³¹), ensure each agent gets to contribute, and prevent chaos or infinite loops in the discussion ³². It can also inject system messages or instructions to guide the agents if they stray off-topic or get stuck, essentially acting as a moderator. In some cases, the "manager" role can be partially filled by a human operator supervising the chat, especially if the discussion needs steering toward productive outcomes ³³. The design should allow for **human-in-the-loop** oversight, meaning a human can intervene or take over the chat manager role

when needed ³⁴ – for example, to stop an unproductive debate or to provide external information to the agents.

An important aspect to handle is **shared context and memory**. As agents collaborate, the system must maintain a coherent conversation state that accumulates all relevant points raised so far ³⁵. This context (or a summarized form) should be fed into each agent's prompt so they are aware of what others have said. In long debates with multiple rounds, having a **context persistence layer** or memory buffer is vital so that earlier arguments aren't forgotten ³⁶. For instance, *DebateSim* – a research prototype for legislative debates – included a memory manager to enforce cross-round coherence, ensuring the AI debaters remembered and addressed prior points across five rounds of argument ³⁶. Our system might use a similar strategy: e.g., maintain a transcript of the dialogue or a running summary that gets appended to each agent's input each turn. This helps prevent agents from repeating points or contradicting themselves unknowingly.

Termination criteria need to be defined as well. Unlike a fixed pipeline which ends after the last stage, a free-form multi-agent discussion could potentially go on indefinitely. The system should detect or decide when a “good enough” solution or consensus has been reached, or when to stop the debate. This could be a simple fixed number of turns/rounds after which the conversation is halted and results aggregated ³⁷, or a more dynamic check (e.g., agents converge on the same answer, or a moderator/judge agent determines the debate has run its course). Designing robust decision-making mechanisms for **when and how the agents output a final answer** is crucial ³⁸. Some systems employ a dedicated *judge* or *arbiter* agent for this purpose – either an AI agent tasked with evaluating arguments and picking a winner, or a voting mechanism among the agents themselves. For example, the MIT CSAIL multi-model collaboration strategy has all agents critique each other's answers through several rounds, then uses a **majority vote** across the agents' final answers to produce the unified output ². Similarly, Microsoft's *Autogen* framework demonstrates a debate pattern where multiple solver agents exchange solutions, and an aggregator agent collects their final responses to decide the answer by majority voting ¹⁸. These approaches ensure the outcome reflects the collective reasoning of the group.

Example Use Case: AI Debate Simulation

One envisioned use case is a **debate between AI “government ministers”** – essentially a simulation where each agent represents a different minister or policy perspective (e.g. an AI Minister of Finance, Minister of Environment, Minister of Defense, etc.) and they debate a policy proposal. This scenario maps well to the *group chat orchestration* pattern. Each minister agent would have a role-specialized profile (different knowledge and priorities) and the system would initiate a structured dialogue on a given topic (say, a new climate policy or budget plan). To make this work, the **coordination layer** might assign initial speaking turns or questions to each agent, and enforce that the agents respond in a logical order (perhaps mimicking parliamentary debate rules). The agents would present arguments for or against aspects of the proposal, challenge each other's points, and attempt to reach a resolution or recommendation.

Research on multi-agent debate provides guidance for such a setup. *DebateSim*, for instance, implemented a legislative debate with **Pro and Con agents** arguing opposite sides of a bill, along with an **AI judge** to provide feedback and evaluate arguments ³⁶. Each agent had to **cite evidence** (integrating with a database of legislative documents) and follow a structured format of opening statements, rebuttals, and closing statements across multiple rounds ³⁶. The presence of an impartial judge agent (or it could be a moderator agent in our case) ensured the debate stayed on track and that there was a mechanism to decide which side's points were stronger in the end. Even if our “AI ministers” scenario is more collaborative than adversarial (since presumably all ministers share a goal of good

governance, though with differing priorities), a similar structure can be used: agents articulate their viewpoints and constraints, and a concluding phase either merges these perspectives into a consensus or has a higher-level agent summarize the final decision.

Another relevant strategy is the **AI-vs-AI debate for truth-seeking**. OpenAI and others have proposed that having two AI agents argue opposite sides of a question can help surface the truth, with a human or a separate model judging which argument is more convincing (the *AI debate* technique) ³⁹. In our system, we could simulate a contentious issue by assigning two agents opposing stances (e.g., one minister advocating a policy, another questioning it) and let them debate to stress-test the policy. This *adversarial collaboration* can highlight pros and cons for the human observer or a judging mechanism. Notably, recent work found that when multiple models engage in such debate with rounds of critique, the **answers tend to become more factual and consistent**, as the process “forces” each agent to justify itself and correct errors that others point out ² ⁴⁰. By the end of a debate, the system could either present the synthesized consensus or both sides’ final stances, depending on the goal (consensus vs. showing both viewpoints).

Design considerations for debate simulations: We should enforce a dialogue protocol (e.g., each agent has a fixed time/turn to speak) and use structured prompts to encourage agents to rebut specific points from others rather than talking past each other. It may help to have an explicit *debate template* the agents follow (for example: “State your argument with evidence -> respond to the previous point -> ask a question to another agent”). Ensuring **grounding in facts** is important – connecting agents to knowledge bases or tool use (like web search or documents) can prevent the debate from devolving into pure opinion. As seen in DebateSim, requiring agents to cite sources led to **89% citation accuracy** and more evidence-grounded arguments in their experiments ³⁶ ⁴¹. Our system might integrate retrieval tools so that each minister agent can pull up relevant data (budget numbers, environmental reports, etc.) to support their statements. Finally, decide how to conclude: perhaps the Prime Minister agent (or a designated moderator) listens to all arguments and then formulates a final policy decision, or the system triggers a voting among the agents on different options.

Example Use Case: Virtual CTO and Advisory Panel

Another scenario is a “**virtual CTO’s office**” where an AI acting as a Chief Technology Officer consults a panel of AI advisor agents on some issue (e.g. deciding on an IT strategy, solving a technical problem, or evaluating a new product idea). This is essentially a **multidisciplinary meeting** in which each advisor agent has a distinct area of expertise – for example, one could be a “Finance Advisor” focusing on budget impact, another a “Security Advisor” focusing on cybersecurity risks, another an “Engineering Lead” focusing on technical feasibility, and so on. The CTO agent (or the system user acting as CTO) poses a question or problem, and the advisor agents discuss it among themselves and with the CTO to provide a well-rounded analysis.

This use case fits the *group chat orchestration* pattern as well, emphasizing **collaborative brainstorming and review**. The system design would have a set of advisor agents who all receive the context of the CTO’s question. The conversation could be less formally turn-based than the debate scenario; instead it might be more of an open discussion, but guided by the chat manager to ensure each advisor contributes. For instance, the CTO (human or AI) could start by asking, “What are the potential pitfalls of adopting Technology X in our stack?” Immediately, the Security Advisor agent might chime in about security compliance issues, the Finance Advisor about cost implications, and the Engineering Lead about integration complexity. They can **build on each other’s points**, creating a rich discussion that covers multiple angles. This is akin to how the **city parks proposal example** from Microsoft’s guidance worked: the system brought in an environmental planning agent, a community

engagement agent, and a budget/operations agent to *debate different community impact perspectives and work toward consensus* on a development proposal ²⁶. Each agent added their expert viewpoint (environmental impact, community feedback, cost and maintenance), and through the managed discussion they identified concerns and synergies, leading to a more refined proposal ⁴² ⁴³. In our virtual CTO meeting, similarly, the outcome would be a well-considered decision or a set of recommendations that factor in all the expert insights.

To implement this, we should ensure that each advisor agent has a clearly defined domain knowledge (perhaps via prompt engineering that sets its persona and scope). The **prompting for each agent** might include a description like “You are the Chief Security Officer, focus on security and privacy issues.” This keeps their contributions focused and complementary. The chat manager can either let the conversation flow naturally or specifically prompt agents in sequence (“Let’s hear from Finance next, what do you think?”) depending on how structured we want the meeting to be. Because this is a collaborative (not adversarial) setting, we might also incorporate a **consensus phase**: after discussion, the CTO agent (or the system) asks for a summary or recommendation that *everyone agrees on* (or each advisor signs off on). Alternatively, the CTO agent could be the one to synthesize the advisors’ inputs and make a decision – similar to how a human CTO would listen to the team then conclude. The **MetaGPT framework** provides an inspiration here: it essentially creates an “AI software company” in which different agents assume roles like Product Manager, Architect, Engineer, etc., and through a coordinated workflow they produce outputs like design documents and code ⁴⁴ ³⁰. MetaGPT’s success demonstrates that assigning **distinct roles to GPT agents and orchestrating their communication** can effectively emulate a team of specialists tackling a project ³⁰. Our system can adopt a similar multi-role strategy for the CTO scenario, ensuring that each advisor agent’s output (ideas, warnings, analyses) is merged into the final plan.

One might consider a slight variation: a **maker-checker loop** within this advisor setting. For example, the CTO could ask one agent (say the Engineering Lead) to draft a proposal or plan, then ask another agent (say a Risk Officer) to **critique or “check”** that draft ³¹. The critique might lead to revisions, and the loop can repeat – this is basically a structured two-agent review cycle embedded in the larger group conversation. Such iterative refinement (maker vs checker) ensures quality control and could be a pattern the system supports for certain tasks (code review, policy review, etc.) ⁴⁵.

Ensuring Flexibility and Robustness

To support all the above use cases, the architecture should be highly **configurable**. It needs a core infrastructure for spawning agents with given personas, managing message passing between them, and integrating results – but the specific conversation pattern might be configured per session or per task. For example, the system could have a mode for “Debate Simulation” (with predefined roles of proponents, opponents, judges) and another for “Advisory Panel” (with a leader and multiple peer advisors). Under the hood, many components can be shared: the chat manager or controller logic, the context handling subsystem, etc., but with different settings (turn-taking rules, termination conditions, role prompts). Leveraging existing frameworks can accelerate development; tools like Microsoft’s **AutoGen** or libraries like Camel and Langroid provide abstractions for multi-agent conversations ⁴⁶ ⁴⁷. These frameworks implement patterns such as debating agents or cooperative agents and handle some of the prompt orchestration and message routing behind the scenes. For instance, AutoGen’s *GroupChat* and *Debate* patterns mirror what we’ve described, and could be extended to our specific roles easily ⁴⁸.

Scalability is another consideration. We should design for a manageable number of agents initially – perhaps 2 to 5 agents in a conversation. As noted in guidance, adding too many agents can make it

hard to control the dialogue and can increase noise or repetition ^{32 49}. Empirically, even three well-chosen agents can significantly enrich a discussion (as seen in the parks department example with three agents ⁴²). If we do need to scale to larger “crowds” of agents for certain applications, we might need hierarchical structures (e.g., clusters of agents each producing a sub-result, then a higher-level agent integrates those – akin to dividing a big meeting into breakout sessions). The literature indicates that performance doesn’t monotonically increase with more agents; beyond a point it can saturate or even worsen due to coordination overhead and conflicting signals ⁵⁰. Our system should therefore allow *scaling out* to more agents or more debate rounds, but also provide mechanisms to mitigate the downsides (like limiting unnecessary communication, grouping agents, or smarter aggregation techniques).

Finally, we must handle the **computational cost** and latency implications. Running multiple large models in parallel or in extended debates is resource-intensive. Caching intermediate results or using smaller specialized models for certain roles could be ways to optimize. One promising angle is that not all agents need be large models – perhaps some roles (like a simple fact-checker) could be a lightweight model or a rules-based system. The architecture could mix LLM-based agents with tool-based or programmatic agents (for example, an agent could be a Python function that checks math, while another is an LLM generating language content). This hybrid approach can save time and ensure accuracy in areas where deterministic tools excel. Nevertheless, when multiple LLMs *are* needed, the system might run them in parallel where possible to reduce total turnaround time (using the concurrent pattern), and only use sequential or iterative patterns when necessary for quality.

In summary, a contemporary multi-agent AI system should **support a spectrum of interaction patterns** – from linear pipelines to free-form group debates – to accommodate use cases like AI debates and virtual advisory panels. Through careful orchestration (via a chat manager or similar controller) and well-defined agent roles, such a system can harness the “wisdom of crowds” effect among AI models. The design must manage the additional complexity: maintaining shared context, preventing conversational dead-ends, and deciding how to consolidate the agents’ contributions. If implemented well, the result is a powerful platform where, for example, *AI government ministers* can debate policies or a *virtual CTO* can consult expert AIs, and the outcome will be a **more nuanced and vetted decision** than any single model would produce on its own. This aligns with findings that multi-agent approaches, while more complex, **yield superior results on challenging problems** – provided we handle their unique challenges responsibly ⁶. The potential applications are wide-ranging, and our system’s flexibility will allow it to support many patterns of AI collaboration as these use cases continue to expand.

Sources:

- Microsoft Azure Architecture Center – *AI agent orchestration patterns* ^{9 15 22 5 26}
- MIT CSAIL (2023) – Multi-model collaboration research ^{2 40}
- *DebateSim* (2025) – Multi-agent legislative debate architecture ^{36 41}
- MetaGPT – Multi-agent software team framework (2023) ^{44 30}
- Tillmann, *Literature Review of MA-LLMs for Problem-Solving* (2025) ^{6 1 51}

^{1 3 4 6 38 50 51} Literature Review Of Multi-Agent Debate For Problem-Solving
<https://arxiv.org/html/2506.00066v1>

^{2 40} Multi-AI collaboration helps reasoning and factual accuracy in large language models | MIT News | Massachusetts Institute of Technology
<https://news.mit.edu/2023/multi-ai-collaboration-helps-reasoning-factual-accuracy-language-models-0918>

5 7 8 9 10 11 12 13 14 15 16 17 19 20 21 22 23 24 25 26 27 28 31 32 33 34 35 42 43

45 49 **AI Agent Orchestration Patterns - Azure Architecture Center | Microsoft Learn**

<https://learn.microsoft.com/en-us/azure/architecture/ai-ml/guide/ai-agent-design-patterns>

18 37 47 48 **Multi-Agent Debate — AutoGen**

<https://microsoft.github.io/autogen/stable//user-guide/core-user-guide/design-patterns/multi-agent-debate.html>

29 30 44 **MetaGPT: A Multi-Agent Framework Revolutionizing Software Development | by Alexei Korol | Medium**

<https://medium.com/@korolalexei/metagpt-a-multi-agent-framework-revolutionizing-software-development-f585fe1aa950>

36 41 **DebateSim: CoT Drift in Multi-Agent Debate Systems in an Architectural and Empirical Study | OpenReview**

<https://openreview.net/forum?id=5yxNX2pmhJ>

39 **Debate Multi-Agent Architectures - Swarms**

https://docs.swarms.world/en/latest/swarms/structs/orchestration_methods/

46 **Exploring 2 Multi-Agent LLM Libraries: Camel & Langroid | blog_posts – Weights & Biases**

https://wandb.ai/vincenttu/blog_posts/reports/Exploring-2-Multi-Agent-LLM-Libraries-Camel-Langroid--Vmlldzo1MzAyODM5